Project 1 FYS-STK4155 Autumn 2017

Jon Audun Baar & Mikael Ravndal

6. oktober 2018

All of the plots of the models are done with degree equals 5.

Introduction

In this project we were to analyse different sets of data; test-data from the Franke-function and real data, using different linear regression models. The main goal of the project was to evaulate the different models, and come to a conclusion in terms of which models would fit best with which data.
Most of our code is in the GitHub repository. We have chosen to not use the code in our report.
We have, however, made a python file, project01.py, that provides most of our code in sequence, so that it is possible to follow the code that produce our tests and plots in order.

Ordinary least square on test data

First we have generated some test data, with the noise being very little, which we have plotted to get a more intuitive feel for how it looks. We have done this for all the models.
Ordinary Least square of the Frankefunction(the same as Ridge with $\lambda = 0$):

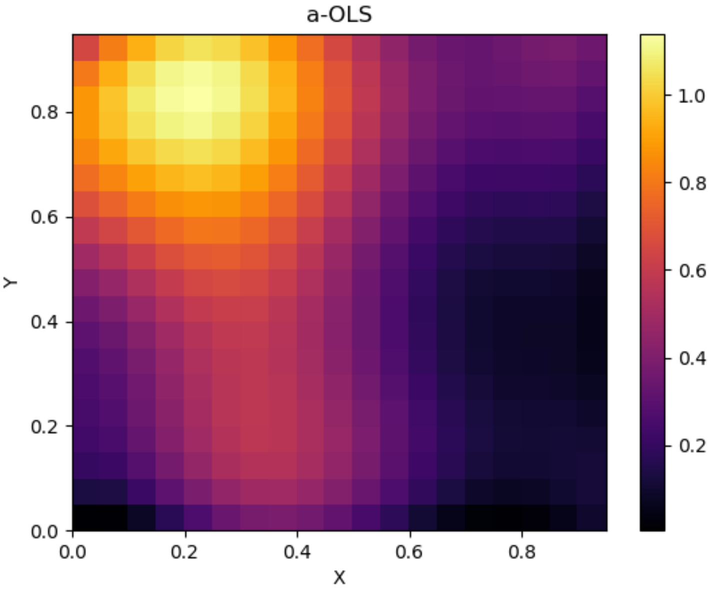

As you can tell from our code, we haven't made an own function for ordinary least square, since it is the same as Ridge, just with the λ set to 0. This will also show later that when the λ gets low it is very similar to OLS.
Here is the values for calculating the five first degrees and their MSE and R2-score:
OLS Test Data We can tell that our model is fitting our test data better and better with a higher degree. Now we have to check with bootstrap as well to see if our model fits the data good.

k	MSE	R2
1	0.020836568820240136	0.7159598591230791
2	0.015631136506458913	0.7869193218104034
3	0.007175355204257789	0.9021869233537347
4	0.003960217232650187	0.9460150723293508
5	0.0018602479352178107	0.9746414541595732

Before the resampling we have also calculated the betas of $k = 5$:
Var of Beta, degree 5

$$\begin{array}{rrrr}
[9.47238055e-03 & 8.01217226e-01 & 4.12183842e-01 & 9.89887798e+00 \\
8.39455751e+00 & 4.14211294e+00 & 2.71771303e+01 & 3.28888083e+01 \\
1.78285914e+01 & 1.31093669e+01 & 1.91368860e+01 & 2.67004422e+01 \\
1.91391193e+01 & 9.34737951e+00 & 1.13215581e+01 & 2.55450522e+00 \\
4.40599813e+00 & 3.60735967e+00 & 1.80107762e+00 & 1.00979408e+00 \\
 & & & 1.50021893e+00]
\end{array}$$

Also the 95-percentage confidence interval of the betas: 95-percentage CI of betas, degree 5

$$\begin{array}{rr}
[[7.25873511e-02 & 4.54098870e-01] \\
[7.23335269e+00 & 1.07421092e+01] \\
[3.44378188e+00 & 5.96043621e+00] \\
[-4.41211552e+01 & -3.17880887e+01] \\
[-2.42828028e+01 & -1.29254533e+01] \\
[-1.57022805e+01 & -7.72437196e+00] \\
[4.22571363e+01 & 6.26923830e+01] \\
[4.18269196e+01 & 6.43072224e+01] \\
[1.67482166e+01 & 3.32996879e+01] \\
[-1.02120404e+01 & 3.98078784e+00] \\
[-3.30453851e+01 & -1.58973753e+01] \\
[-7.33667324e+01 & -5.31114961e+01] \\
[-1.94596236e+01 & -2.31061330e+00] \\
[-3.98768180e+01 & -2.78922323e+01] \\
[1.90552365e+01 & 3.22448231e+01] \\
[-2.43921072e+00 & 3.82593943e+00] \\
[1.88142647e+01 & 2.70423776e+01] \\
[8.68290825e+00 & 1.61280472e+01] \\
[-7.84195698e+00 & -2.58124771e+00] \\
[1.66275865e+01 & 2.05666637e+01] \\
[-1.80168037e+01 & -1.32155417e+01]]
\end{array}$$

The code and commenting for the calculations is to be found in python-file project01.py

Resampling

Using our bootstrapping algorithm with a resampling of 100, degree of five, we get these values:
VAR: 0.000052
BIAS: 0.001933
Bootstrap mean of MSE: 0.0020
Bootstrap mean of r2Score: 0.9757

The bootsrap values aligns pretty well with our original ones.

Ridge regression

Ridge Regression with $\lambda = 0.1$
Graphic plot of how it looks:

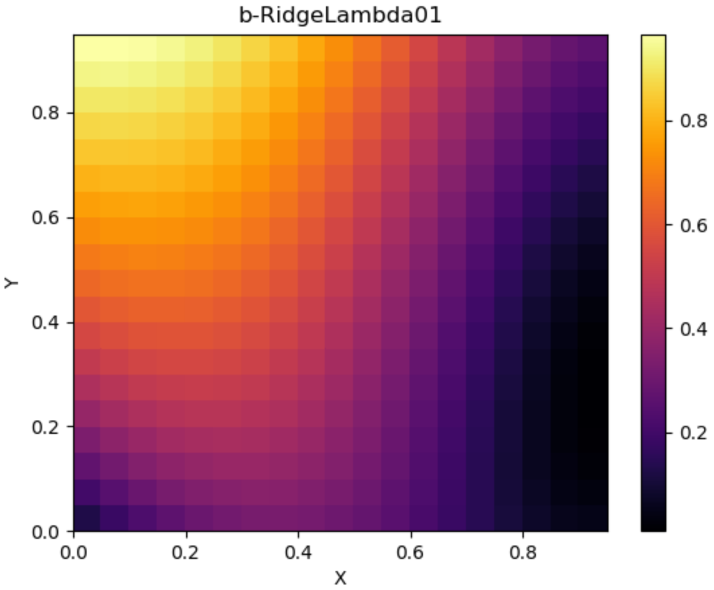

Ridge Test Data

k	MSE	R2
1	0.025965982389446095	0.6906471643146342
2	0.018247163430545398	0.7826074259086047
3	0.010258352759015437	0.8777843076427536
4	0.009382588378732645	0.8882179662029949
5	0.009143926633340667	0.8910613282063765

Compared to OLS, we can tell that Ridge does significantly worse then OLS.
Var of Beta, degree 5

```
[0.00077004  0.01729839  0.01384995  0.06816008  0.06826285  0.05696662
  0.0417508  0.02034709  0.06671475  0.03500783  0.02294549  0.01073849
 0.03470761  0.01261125  0.02127092  0.03349824  0.01062464  0.03548007
                                   0.03422875  0.04204926  0.02520149]
```

Also the 95-percentage confidence interval of the betas: 95-percentage CI of betas, degree 5

$$\begin{array}{rr}
[[8.39518095e-01 & 9.45603016e-01] \\
[6.00189901e-01 & 1.10406786e+00] \\
[9.22637923e-01 & 1.52633094e+00] \\
[-6.51421032e+00 & -5.30387095e+00] \\
[5.84279933e-01 & 1.67511842e+00] \\
[-5.82384066e+00 & -4.85326698e+00] \\
[3.42317950e+00 & 4.16050815e+00] \\
[1.31502300e+00 & 2.10193829e+00] \\
[-2.00335018e+00 & -1.07460992e+00] \\
[1.38884378e+00 & 1.92131292e+00] \\
[3.21253289e+00 & 3.98989176e+00] \\
[5.79118852e-01 & 1.17557901e+00] \\
[-3.41043372e-02 & 7.77207310e-01] \\
[-8.02285903e-01 & -2.23570539e-01] \\
[3.31692208e+00 & 3.80174824e+00] \\
[-3.69457810e+00 & -2.99905374e+00] \\
[-2.10226995e+00 & -1.54360569e+00] \\
[-5.64130918e-03 & 8.94344017e-01] \\
[-1.07674065e+00 & -1.87377295e-01] \\
[3.14047856e-01 & 8.79022520e-01] \\
[-1.87908328e+00 & -1.35609332e+00]]
\end{array}$$

Resampling

We can take a look at how different lambdas and different degrees of the polynomial makes a change in the R2-score and the MSE.
Here is a plot to show how they develop as a function of lambda.

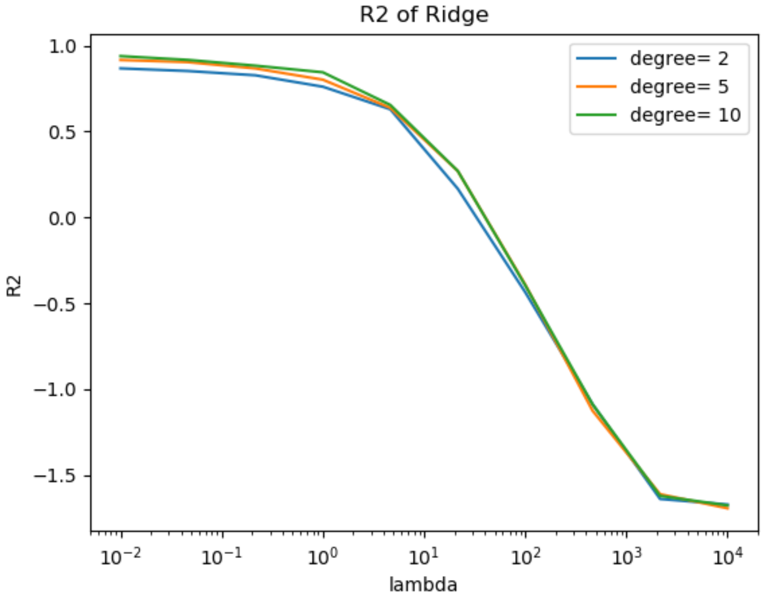

We can tell pretty easily that the degree of the predictions doesn't matter much compared to how much the choice of lambda do. We can still tell that a lower degree function does worse then the other functions.

Some interesting values from bootstrap:

Bootstrap-values from degree of 5, lmb = 0.1 and 100 bootstrap-samples
VAR: 0.000067
BIAS: 0.008640
Bootstrap mean of MSE: 0.0087
Bootstrap mean of r2Score: 0.8980

Bootstrap-values from degree of 5, lmb = 1 and 100 bootstrap-samples
VAR: 0.000059
BIAS: 0.011915
Bootstrap mean of MSE: 0.0120
Bootstrap mean of r2Score: 0.8597

Bootstrap-values from degree of 5, lmb = 10 and 100 bootstrap-samples
VAR: 0.000066
BIAS: 0.020274
Bootstrap mean of MSE: 0.0203
Bootstrap mean of r2Score: 0.7617

Bootstrap-values from degree of 2, lmb = 10 and 100 bootstrap-samples
VAR: 0.000057
BIAS: 0.022754
Bootstrap mean of MSE: 0.0228
Bootstrap mean of r2Score: 0.7327

Part c)

Lasso Regression with $\lambda = 0.01$

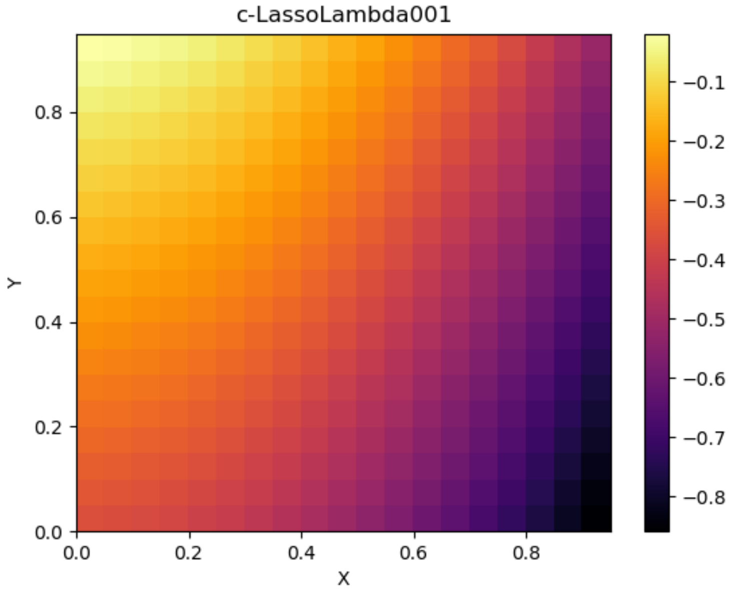

Ridge Test Data

k	MSE	R2
1	0.25272407945476655	-2.01089746779934852
2	0.25272407945476655	-2.0108974677993485
3	0.25272407945476655	-2.0108974677993485
4	0.25272407945476655	-2.0108974677993485
5	0.25272407945476655	-2.0108974677993485

Var of Beta

```
[0.   0.00025829   0.00394575   0.   0.   0.00324571
                     0.   0.   0.   0.   0.   0.   0.
                     0.   0.   0.   0.   0.   0.   0.
                                                   0.]
```

95-percentage CI of betas

$$\begin{array}{r}
[[0. \quad 0.]\\
[-0.40937802 \quad -0.34637908]\\
[-0.17825992 \quad 0.06797132]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0.58986796 \quad -0.36654512]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]\\
[-0. \quad 0.]]
\end{array}$$

There is obviusly something happening with Lasso that doesn't work. Our Lasso method works great with the real data.

Resampling

Bootstrap-values from degree of 1, lmb = 0.1 and 100 bootstrap-samples
VAR: 0.000000
BIAS: 0.239704
Bootstrap mean of MSE: 0.2397
Bootstrap mean of r2Score: -2.0278
This was the same for all degrees and values of lambda, so I only included this.

Part d)

Imports 100x100 chunk of real data from top left corner of dataset nr.1.
Plot of real data

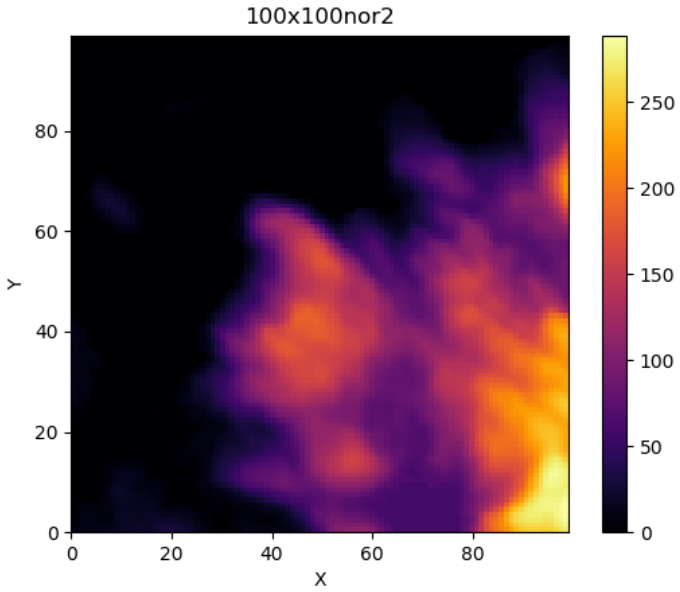

Real Data

Repeat of the previous method and data, but with real data.

OLS

Plot of the data with the OLS-method of degree 5:

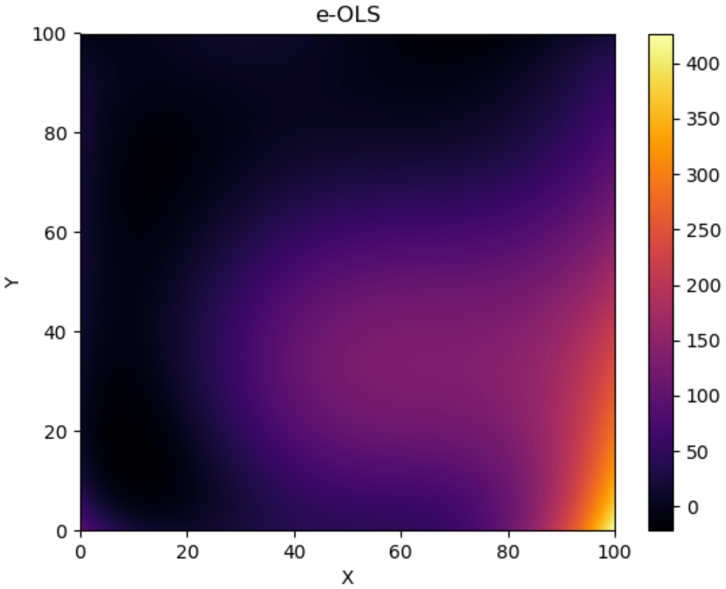

OLS Score of the Real Data:

k	MSE	R2
1	1912.190996	0.587579
2	1177.397306	0.746059
3	1023.042624	0.779351
4	824.587589	0.822153
5	791.268452	0.829340

Var of Beta

$$\begin{array}{rrrr}
[2.32350868e-17 & 3.59021583e-14 & 9.26320612e-14 & 6.67623760e-12 \\
2.76179532e-12 & 2.44594128e-11 & 1.00899287e-09 & 2.82477031e-08 \\
3.26102575e-08 & 4.03014136e-09 & 7.93068781e-13 & 8.29815731e-12 \\
2.72417498e-12 & 8.99036913e-12 & 2.47909085e-12 & 4.07929803e-17 \\
2.67723887e-16 & 5.18767547e-16 & 4.23509770e-16 & 1.84571052e-16 \\
 & & & 9.42227510e-17]
\end{array}$$

Resampling

Bootstrap-values from degree of 5 and 100 bootstrap-samples
VAR: 1.287057

BIAS: 791.273865
Bootstrap mean of MSE: 792.5609
Bootstrap mean of r2Score: 0.8291

The MSE and R2-score are really close to our values for MSE and R2.

Ridge

Plot of the data with the Ridge-method of degree 5, lambda = 0.1:

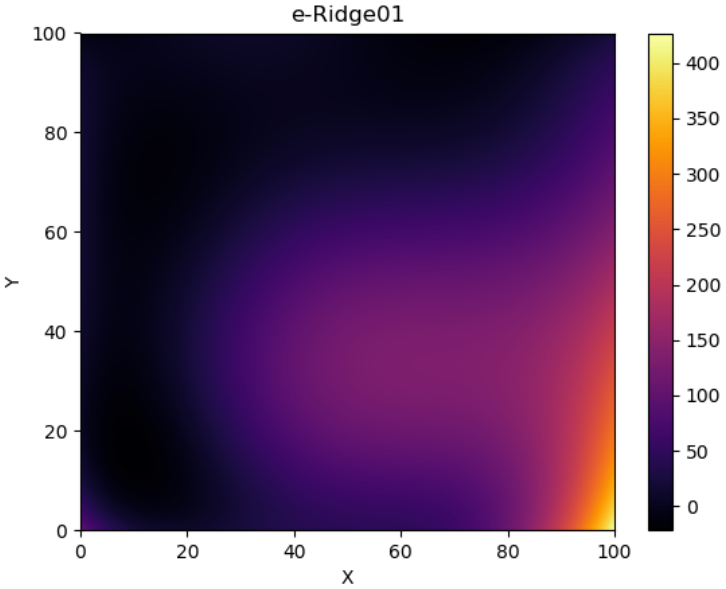

Ridge Score of the Real Data with lambda = 0.1:

k	MSE	R2
1	1912.190999	0.587579
2	1177.397307	0.746059
3	1023.042624	0.779351
4	824.635824	0.822143
5	791.268452	0.829340

Resampling

Bootstrap-values from degree of 5, lmb = 10 and 100 bootstrap-samples
VAR: 1.297002
BIAS: 791.285081
Bootstrap mean of MSE: 792.5821
Bootstrap mean of r2Score: 0.8291
Can tell that MSE and R2-score is pretty similar to our computations.

Lasso

Plot of the data with the Lasso-method of degree 5, lambda = 0.01:

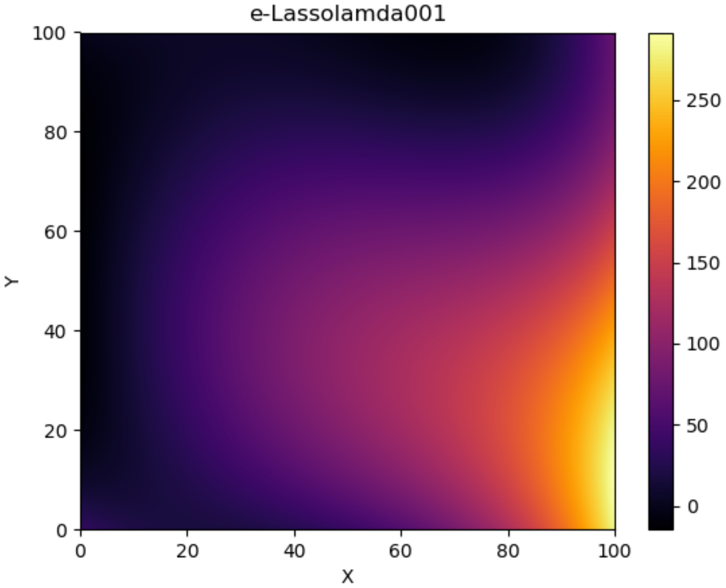

Lasso Score of the Real Data with lambda= 0.1:

k	MSE	R2
1	6601.647721	-0.423841
2	1564.545187	0.662560
3	1189.660192	0.743415
4	1155.338990	0.750817
5	974.391256	0.789844

Resampling

Bootstrap-values from degree of 5, lmb = 10 and 100 bootstrap-samples
VAR: 1.689048
BIAS: 1239.400793
Bootstrap mean of MSE: 1241.0898
Bootstrap mean of r2Score: 0.7323

Here the MSE and R2score is a pretty long way from our scores when we use all our data. So the model is probably a little overfitted and we should expect our R2 to actually be lower than it is.

Conclusion